

Prediction and Choosing Predictors

Unit 4: Predicting a New Case, Overfitting, and Which Variables to Keep

Stat 220 | Statistics for Data Science

Where we left off

- Unit 3 fitted a model and read one coefficient at a time.
- This unit uses the whole model at once, for two jobs: predicting a case you have not seen, and deciding which variables belong in the model.
- Both come down to one question: how will this model do on data it was not fitted on?

What this unit gives you

Things to know

- how to produce a prediction and an honest interval around it
- the difference between an interval for the line and an interval for one new case
- why error measured on your own data is too small
- train/test splits and cross-validation
- what underfitting and overfitting look like
- how to compare candidate sets of predictors

Mistakes to stop making

- quoting a prediction with no interval
- quoting the interval for the line when someone asked about one case
- reporting the fit on the rows you used to fit it
- adding predictors because R^2 went up
- searching many models and reporting the winner as if you had picked it in advance
- predicting outside the range of the data

The data for this unit

`moving_jobs.csv`: 600 completed jobs at a moving company.

Column	What it is
hours	crew time on site, the outcome we predict
volume_cuft	how much is being moved
crew_size	movers assigned, 2 to 4
stairs_flights, miles	flights of stairs, one-way distance
packing_service	1 if the customer bought packing
est_boxes,	also recorded
dispatcher_rating,	
weekend, truck_id	

The company quotes a customer before the job. Quote too low and the crew runs over. Quote too high and the customer books someone else.

A prediction is the fitted equation, evaluated at a new case

$$\hat{y}_{\text{new}} = b_0 + b_1 x_{1,\text{new}} + \cdots + b_k x_{k,\text{new}}$$

- Put the new case's values into the equation you already fitted.
- In code that is one line: `fit.predict(new)`, where `new` is a data frame with the same columns the model was fitted with.
- The number that comes back is a *point* prediction. On its own it is not an answer.

Example: quoting a job

A customer has 800 cubic feet, a crew of 3, two flights of stairs, 12 miles, and wants packing.

predicted hours	12.39
-----------------	-------

- That is the model's best single guess for this job.
- The dispatcher's actual question is different: *how long could this job run?*
- Two customers with identical numbers do not take identical time. The model knows that, and it can tell you how much they vary.

Two different questions, two different intervals

- **Interval for the line.** Where does the average outcome sit, for jobs like this one? It shrinks as the sample grows, because it is about an estimate.
- **Interval for one new case.** Where will *this job* land? It includes the job-to-job variation that no model removes, so it stays wide no matter how much data you collect.

Principle: which one to quote

If the question is about an average, use the interval for the line. If the question is about one case, use the interval for a new case. A dispatcher scheduling one crew is asking about one case.

The same job, both intervals


	95% interval	width
average job like this one	12.25 to 12.54	0.3 hours
this one job	10.74 to 14.04	3.3 hours

- Same model, same customer, same prediction of 12.39 hours.
- The first interval is eleven times narrower. Quoting it to a dispatcher would promise a precision the model never had.

Trap: the default output

Software often prints the narrow one first. Check which interval you are looking at before it reaches a customer.

The two intervals, drawn



The dark band is where the line sits. The pale band is where a single new job lands, and about 95% of the completed jobs fall inside it.

Your turn #1

Your turn: which interval does each question need?


- ① “If we did a hundred jobs like this one, what would our average time be?”
- ② “The Thompson move is Tuesday. How many hours should I block on the schedule?”
- ③ “Across all 600 jobs, is packing service associated with longer jobs?”

Decide with a neighbor which of the three needs the interval for one new case.

Your turn #1: answer

- **(1) The line.** It asks about an average over many jobs.
- **(2) One new case.** One move, one crew, one Tuesday. This is the wide interval, and the one people most often get wrong.
- **(3) Neither.** That is a question about a coefficient, which is Unit 3. It asks what the model says about a relationship, not what will happen next.

Predicting outside the range of the data



- The line was fitted on jobs from 300 to 900 cubic feet.
- Asked about a 1,400-cubic-foot job, it answers without complaint.
- The jobs that size, when they came in, took longer than the line said. Nothing in the fitted data could have told you that.

Error measured on your own data is too small

- The coefficients were chosen to make the residuals small *on those rows*.
- So the error you compute on them is not what you will see on a new job. It is the best case, not the typical case.
- The fix is to measure on rows the model has not seen.

Principle: the only honest number

Report error on data the model was not fitted on. Every other number is a description of how well it memorized.

The same model, both ways

Five predictors, 600 jobs, 70% used for fitting and 30% held back:


	typical error
on the rows it was fitted on	0.79 hours
on the rows held back	0.93 hours

- The gap here is modest, because this model is simple and the sample is large.
- The gap grows with model flexibility, and it is largest exactly when you are most pleased with your fit.

Holding rows back

- **Train/test split.** Fit on most of the rows, measure on the rest.
`train_test_split(X, y, test_size=0.3)`.
- One split is a single draw, so the number wobbles depending on which rows landed where.
- **Cross-validation** averages over several splits, so every row is held out once.

K -fold cross-validation



- Split the rows into K pieces, here 5. Fit on four, measure on the fifth, and rotate.
- Average the five errors. That average is your estimate of error on a new job.
- `cross_val_score(model, X, y, cv=KFold(5, shuffle=True, random_state=0))`

Two ways to get this wrong

- **Using the held-out rows to choose.** Try twenty models, keep the one with the best held-out error, and report that error. It is now a best-of-twenty, which is optimistic for the same reason as before.
- **Touching the held-out rows early.** Filling in missing values, scaling, or picking predictors using all the data leaks the held-out rows into the fitting.


Trap: the number you report

Choosing with held-out data is fine. Reporting that same number as your final answer is not. If the choice mattered, keep a slice you never touched, or say plainly that the number came from the same data you chose with.

Two ways a model can be wrong


- **Underfitting.** The model is too rigid to capture the pattern that is there. It is wrong in the same direction for whole stretches of the data.
- **Overfitting.** The model is flexible enough to follow the noise. It fits the rows it saw and misses the next ones.
- Both show up as large error on new data, which is why you measure there.

The same 40 jobs, three models



Left: a straight line misses the bend at the high end. Middle: one curve, about right. Right: degree 14 passes near almost every point, and the wiggles are noise.

Training error and held-out error, as complexity grows



- Error on the fitted rows falls as you add flexibility, and keeps falling.
- Error on held-out rows falls, bottoms out, then climbs. Here it runs off the top of the chart.
- The bottom of the red curve is the complexity you want.

How much complexity you can afford

When this is true	You can support
more rows	more flexibility
less noise in the outcome	more flexibility
a strong, simple relationship	less flexibility, and you do not need more
few rows, noisy outcome	a simple model, and honest error bars

There is no complexity that is right in general. It depends on how much data you have and how noisy it is, which is why you check rather than guess.

Your turn #2

Your turn: diagnose the model

A colleague reports: $R^2 = 0.98$ **on the data they fitted**, and a typical error of **4.1 hours** on jobs from the following month. The model has 60 predictors and was fitted on 200 jobs.

Is this underfitting or overfitting, and what would you change first?

Your turn #2: answer

- **Overfitting.** Nearly perfect on its own rows, badly off on new ones, with 60 predictors for 200 jobs.
- **First change:** cut the model down to the handful of predictors that have a reason to be there, then compare by cross-validation.
- A useful check: if a model has nearly as many predictors as it has rows, it can fit almost anything, including noise.

Adding a predictor never hurts the fit on your own data

- R^2 cannot go down when you add a variable. Even a column of random numbers nudges it up.
- So R^2 cannot tell you whether a variable belongs. Neither can “the coefficient came out significant,” if you went looking through many variables.
- You need a criterion that charges you for complexity, or a measurement on data you did not fit.

Four numbers people use

Number	What it does
R^2	share of variation explained. Always rises with more predictors, so it cannot choose
adjusted R^2	R^2 with a small charge per predictor. Better, but a weak guard
AIC, BIC	fit with a charge per predictor. Lower is better. BIC charges more, so it picks smaller models
cross-validated error	measured on rows left out. Slower to compute, and closest to the question you care about

AIC and BIC compare models fitted to the same rows with the same outcome. Cross-validated error compares anything that can make a prediction.

Four candidate sets, scored every way

Predictors	R^2	adj. R^2	AIC	BIC	CV error
volume	0.8717	0.8715	1766.7	1775.5	1.056
volume, crew	0.8873	0.8870	1690.9	1704.1	0.991
the five that make sense	0.9197	0.9190	1493.9	1520.2	0.840
those five plus three junk columns	0.9202	0.9191	1496.4	1535.9	0.839

- The junk columns (estimated boxes, dispatcher rating, weekend) raise R^2 , and even nudge adjusted R^2 up.
- AIC and BIC both get worse. Cross-validated error is a tie, 0.001 hours apart.
- Three extra columns bought four seconds of accuracy on a twelve-hour job.

Searching through variables is itself a way to overfit

- Try enough predictors and some will look useful by chance. With 20 useless columns and a 0.05 cutoff, about one shows up “significant.”
- **Stepwise selection** adds or drops variables by p -value, over and over. The model it lands on is the winner of a long search, and its p -values are no longer the p -values of a test you planned.
- The junk columns above had p -values of 0.13, 0.30, and 0.80. That is the pattern to expect from columns unrelated to the outcome, and the reason not to keep hunting until one of them dips under 0.05.

Penalized regression, in one slide

- Ordinary least squares picks coefficients to make the residuals small. **Lasso** adds a charge for the size of the coefficients.
- The effect is that weak predictors shrink toward zero, and some land exactly on zero, which drops them from the model.
- One number, the size of the charge, controls how aggressive it is. You choose it by cross-validation.

In practice: on the moving data

Lasso set weekend to exactly zero and shrank dispatcher_rating to nearly nothing. It kept est_boxes, which is a rough copy of volume, so the two of them split the credit.

A recipe you can defend

- ① **Start from the question.** Predicting, or explaining a particular effect? The answer changes which variables must be in the model.
- ② **Write down a small number of candidate sets**, each with a reason. Three or four, not every subset.
- ③ **Score them the same way**, with cross-validated error, or AIC and BIC if the models are comparable.
- ④ **Take the simplest set that is close to the best.**
- ⑤ **Report how you chose**, including the sets that lost.

Your turn #3

Your turn: is the new column worth keeping?

The company starts recording **elevator (yes/no)**. Adding it moves R^2 from 0.9197 to 0.9206, AIC from 1493.9 to 1492.1, and cross-validated error from 0.840 to 0.835.

Keep it or not? What would you want to know before deciding?

Your turn #3: answer

- Every number moved the right way, and unlike the junk columns, AIC improved rather than got worse. That is mild evidence the column carries real information.
- The gain is 0.005 hours, about 20 seconds. The case for keeping it is that stairs and elevators plainly affect a move, not that the numbers jumped.
- Worth knowing: is it recorded for every job, and is it recorded *before* the quote? A predictor you will not have at prediction time is useless, however well it scores.

LLM check: mistake or not? (1 of 2)

LLM check: we asked an LLM to evaluate a model

We told it: “I fitted a model on all 600 jobs and the R^2 is 0.94. Is it good enough to quote customers with?”

It answered: “An R^2 of 0.94 means the model explains 94% of the variation, so it is performing very well and is suitable for quoting.”

Mistake, or no mistake? Decide before we click.

LLM check: mistake or not? (1 of 2)

LLM check: we asked an LLM to evaluate a model

We told it: “I fitted a model on all 600 jobs and the R^2 is 0.94. Is it good enough to quote customers with?”

It answered: “An R^2 of 0.94 means the model explains 94% of the variation, so it is performing very well and is suitable for quoting.”

Mistake, or no mistake? Decide before we click.

Trap: verdict: mistake

The 0.94 was measured on the rows the model was fitted on, and quoting is a question about one new job. The right answer is a cross-validated error and a prediction interval.

LLM check: mistake or not? (2 of 2)

LLM check: same data, a different answer

We told it: “Adding three more columns raised R^2 from 0.9197 to 0.9202. Should I keep them?”

It answered: “That rise is what R^2 does when you add any column. Check a criterion that charges for complexity, such as AIC, or compare cross-validated error, and keep the columns only if they earn it.”

Mistake, or no mistake?

LLM check: mistake or not? (2 of 2)

LLM check: same data, a different answer

We told it: “Adding three more columns raised R^2 from 0.9197 to 0.9202. Should I keep them?”

It answered: “That rise is what R^2 does when you add any column. Check a criterion that charges for complexity, such as AIC, or compare cross-validated error, and keep the columns only if they earn it.”

Mistake, or no mistake?

Principle: verdict: no mistake

This is the right move, and it names the reason: R^2 rises mechanically, so it cannot be the deciding number.

Producing a prediction you can defend

- ① **Fit the model** on the data you have.
- ② **Check the range.** Is the new case inside the data the model was fitted on?
- ③ **Predict:** `fit.predict(new)`, with the same columns as the fit.
- ④ **Decide which interval** the question needs: the line, or one new case.
- ⑤ **Measure error honestly**, on held-out rows or by cross-validation.
- ⑥ **Report the number with its interval**, and say what it is an interval for.

The traps, all in one place

Trap: the ones that bite most often

- Reporting a prediction with no interval.
- Quoting the interval for the line when the question was about one case.
- Reporting R^2 or error from the rows the model was fitted on.
- Adding predictors because R^2 rose.
- Running a long search and reporting the winner's p -value as if it were planned.
- Choosing a model on held-out data and reporting that same error as the final number.
- Predicting outside the range of the data.
- Keeping a predictor you will not have when the prediction is actually needed.

What you should be able to do with software

You are not expected to memorize syntax. You are expected to know which procedure the question calls for, what to hand it, and what the output means.

You should be able to	What you would reach for
Predict a new case	<code>fit.predict(new_df)</code>
Both intervals at once	<code>fit.get_prediction(new_df)</code> <code>.summary_frame()</code>
Hold rows back	<code>train_test_split(X, y, test_size=0.3)</code>
Cross-validate	<code>cross_val_score(model, X, y,</code> <code>cv=KFold(5, shuffle=True))</code>
Compare models by AIC or BIC	<code>fit.aic</code> , <code>fit.bic</code>
Penalized regression	<code>LassoCV(cv=5).fit(X_scaled, y)</code>

Where we go next

- This unit asked how a model behaves on data it has not seen, and which variables deserve to be in it.
- Unit 5 asks a different question: where the data came from in the first place, and when a coefficient can be read as an effect rather than an association.
- A model can predict well and still be the wrong basis for a decision about what to change. That distinction is the next unit.